

REGIME LAB

A Bayesian Regime Detection Engine for Indian Equity Markets

Zetheta Algorithms | Task 1A | Days 1-14 of 15

156 tests passing | 11 executed notebooks | 18 commits | every claim checked against output before being written

The Question This Engine Answers

- What regime is the Indian equity market in right now, with calibrated uncertainty, not a point guess?
- What should a fund do about it: tilt rules conditioned on regime AND conviction, not just regime alone.
- Can every number be traced back to the model, the data window, and the seed that produced it?

Core design constraint

Point forecasts cannot be defended to an Investment Committee or SEBI. Every output in this engine is a probability distribution with a traceable lineage, not a single number asserted with confidence it has not earned.

Pipeline Architecture: Raw Data to IC Artefact

1. Data

Synthetic generator + 2 real partial datasets

2. Features

31 tabular + TDA + GCN (GCN blocked: torch broken)

3. Regime models

HMM, Bayesian HMM, RS-VAR, BDL x3, foundation head

4. Ensemble

BMA + stacking, PSIS-LOO selection

5. Online layer

Particle filter, BOCPD, streaming updates

6. Calibration

Split-conformal, APS, Mondrian, ACI

7. Decision

Tilt overlay, Monte Carlo VaR/CVaR

8. Output

IC artefact, full lineage

Methodology: Frequentist and Bayesian HMM (Days 4-5)

- 5-state Gaussian HMM, hmmlearn, 20-restart robustness fitting.
- BIC prefers 3 states over the briefed 5; ground-truth match ~40% vs ~20% random baseline.
- Bayesian HMM (PyMC, 4 chains x 500 draws): found the spec's own A3.4 code is not a working HMM (posterior mean equals prior mean), fixed via explicit forward-algorithm marginalisation.
- Cleaner MCMC did not mean better recovery: Bayesian match rate 21% (chance-level) vs frequentist 29%.





Fig: Bayesian HMM regime overlay vs. frequentist, same window.


Methodology: Multivariate Regime-Switching VAR (Day 6)

- 6-D: returns, vol, breadth proxy, INR return, gilt yield change, FII flow proxy.
- Found two real spec bugs: a nonexistent PyMC API call, and LKJCholeskyCov not batching correctly across K regimes.
- Compute-limited to K=3 and 320 posterior draws (4 chains x 80).
- R-hat 1.44-2.43 despite zero divergences: diagnosed as label-switching, fixed by post-hoc relabelling (R-hat drops to 1.00-1.18).
- PSIS-LOO (Day 9): $p_{loo} \sim 205$, close to the ~ 197 raw parameter count - the posterior is genuinely under-sampled, not just noisy.



Methodology: Bayesian Deep Learning (Day 7)


- Three architectures: MC Dropout, variational BNN, deep ensemble (M=10).
- Labels are 88.6% one class (Risk-On): accuracy is uninformative by itself.
- None beats 'always predict Risk-On' on raw accuracy.
- Variational BNN beats the class-prior baseline on NLL in 5 of 5 folds - the most consistent, not the most accurate.
- Uncertainty AUROC 0.84-0.86 for flagging wrong calls; roughly 1 in 5 missed stress days is missed with high confidence.



Methodology: Foundation Models (Day 8)


HuggingFace unreachable from this environment: these are synthetic-pretrained re-implementations of Chronos/TimesFM architectures, not real checkpoints. Stated once, applied everywhere this result is cited.

- Chronos-style (causal transformer) and TimesFM-style (patch decoder) architectures, pretrained on a synthetic corpus.
- Pretraining helps on trend/seasonal structure; near zero gain on return-like series.
- Sample efficiency: pretrained embeddings beat untrained (AUROC 0.875 vs 0.5), but a trivial scale-free vol-ratio baseline beats the embeddings (0.93-0.96 vs 0.875).




Ensembling Design: BMA, Stacking, PSIS-LOO (Day 9)

- 7 members combined via Bayesian Model Averaging and constrained stacking, 4-fold TimeSeriesSplit.
- Central honest finding: the ensemble beats every individual member in only 1 of 4 tested cases.
- Against the true regime, naive equal-weighting beats BOTH BMA and stacking.
- Mechanism, not mystery: too little data for the number of combination weights being fit.
- 3 of 7 members are not genuinely out-of-fold on the evaluation window - flagged prominently, not buried.




Sequential Inference: Particle Filter and BOCPD (Day 10)

- Bootstrap particle filter (5,000 particles, vectorised): mean TV distance 0.0095 vs. exact forward filtering, 99.2% hard-label agreement.
- BOCPD: found two real bugs in the spec's own example. Default priors mismatched to real data by ~4 orders of magnitude.
- $R[0,t]$, the obvious 'changepoint probability' reading, is PROVEN (algebra + numerics) to equal the hazard rate always, regardless of data.
- Fixed: calibrated priors + corrected signal extraction detect real breaks cleanly.




Conformal Prediction and Calibration (Day 11)

- Split-conformal + APS verified over 50 trials (coverage holds on average, not just once).
- Split-conformal produces genuinely EMPTY prediction sets ~10% of the time - confirmed, matching the method's own documented weakness.
- Mondrian conformal: coverage ranged 0.53-1.00, traced to a real, MEASURED temporal shift in one model's own calibration quality, not just small-sample noise.
- ACI: proved a real limit - α_t cannot push $q_{\hat{t}}$ past $\max(\text{calibration scores})$, so severe shift can defeat online adaptation entirely.




Calibration Evidence: Reliability and ECE




Every model's ECE against the TRUE regime (0.30-0.39) is 6-10x its ECE against the HMM training labels (0.03-0.07): a calibration number is only meaningful relative to the label set it was computed against.

Case Study: The 2018 IL&FS Credit Shock (Day 12)





- Real Nifty 50 large-cap data: -14.55% drawdown, closely matching the briefed figure.
- Daily returns around the actual default date: -1.33% to +1.27%, within ~1.5 std of normal volatility.
- Day 10's BOCPD, re-checked here, fires nowhere near September 2018 on large-cap returns alone.
- The strongest evidence IN THIS PROJECT'S OWN RESULTS for why the RS-VAR's cross-asset features matter.

Case Study: BOCPD Tested Fresh on Real VIX (Day 12)



- COVID: calibration window ran directly through the crisis onset (VIX 11.5 to 71.9 within the 60-day window) - the detector's 'normal' baseline was contaminated by the crisis itself.
- Election: fires correctly at a lower threshold; 3 of 5 detected real-data event clusters independently confirmed via web search (Aug 2015, Nov 2016, Sep 2019).

Allocation Overlay Backtest and Monte Carlo (Day 13)



- Point-in-time-valid: refit on pre-2019 data only, not the full series, to avoid look-ahead bias Day 9 flagged elsewhere.
- K=5 and K=3 both fail on 8 years of data; K=2 is what this window actually, robustly supports.
- Post-Shock has the SECOND-HIGHEST drift of any regime (0.25): a naming-based tilt rule gets this exactly backwards.
- Synthetic price path explicitly flagged throughout - the real market did very differently over the real 2019-2024 period.

Cross-Language Validation: Python vs. R (Day 14)

Check	Python	R	Finding
HMM log-lik (K=2)	12073.49 (30 restarts)	12077.19 (10 restarts)	Different local optima - independent confirmation of weak identifiability
Conformal q_hat	0.9934	0.9934 (after fix)	Exact match after fixing a quantile-interpolation bug
VaR(5%)	-0.4900	-0.4942	Agree within expected MC sampling noise

CRAN is unreachable from this sandbox (checked directly): the briefed R packages (depmixS4, MSwM, bcp, conformal, mlr3) are uninstallable. Implemented the underlying algorithms in base R instead - a stronger cross-check than a packaged alternative, since it tests whether two independent implementations of the same mathematics agree.

Regime Lab: The Project IS the Gamification Concept (Part B)

This repository's name is not a coincidence: Part B of the brief describes a gamified training platform called Regime Lab with a 9-level campaign. The technical work built in Days 4-11 maps onto it almost exactly.

Level	Brief's requirement	What this project actually did
3: HMM Engineer	Identify 2008/13/18/20 breaks within 5 days	Day 4: 5-state HMM, ground-truth match ~40% (not evaluated against the literal historical dates)
4: Bayesian Sceptic	Conformal coverage $\geq 88\%$	Day 5 + 11: Bayesian HMM built; split-conformal coverage ~0.90-1.00 on real model data
5: Multivariate Modeller	Regime-conditional covariance w/ credible intervals	Day 6: built, but credible intervals too wide to claim this with confidence
8: Ensembler	Beat every member on OOS log-likelihood	Day 9: did NOT, in 3 of 4 tested cases - reported honestly, not hidden
9: Regulator's Choice	Full lineage, calibration, reconciliation, audit trail	Day 13-14: IC artefact lineage, Day 10 reconciliation, Day 14 peer review

What the Evidence Actually Supports

Regime ID from returns alone is weak

Confirmed 3 independent ways: Day 4's BIC, Day 13's $K=2$ -not- $K=5$ finding, Day 14's Python/R local-optima divergence

The ensemble does not reliably win

1 of 4 tested cases (Day 9); equal-weighting beats BMA and stacking against the true regime

BOCPD needs per-series calibration

Found 3 real events (Day 10); missed 2 of 3 real crises tested in Day 12, each for a specific, traced reason

Calibration drifts over time

Day 11 found a MEASURED, not theoretical, exchangeability violation in this project's own data

Naming traps are real and costly

Post-Shock (Day 13) looks risky by name; it has the 2nd-highest drift of any regime

Where This Leaves the Project

- 14 of 15 days complete: every component built, tested, and documented with its own honest limitations.
- Next priorities, in evidence-based order: more RS-VAR posterior samples, the event-calendar conviction dampener, per-series BOCPD threshold calibration, real continuous data acquisition.
- Deliverable 1 (40-page report): content exists across 14 day-notes files, exceeding 40 pages by volume; a single consolidated PDF remains a mechanical, not a content, gap.
- This deck, the repository, and docs/EXIT_DOCUMENTATION.md are the handover.